



Get the benefits of C without leaving Python

Chris LeBlanc
GNS Science

Myself

- Background in Earth Sciences, Geophysics
- Using Python since 2001
- Software developer for GLOBE Claritas, GNS Science

What is Cython

- Fork of Pyrex
- Easy Python C extensions
- Performance boost
- Python -> C bridge
- C -> Cython bridge

Python Demo

from <https://docs.python.org/2/c-api/intro.html>:

```
def incr_item(dict, key):  
    try:  
        item = dict[key]  
    except KeyError:  
        item = 0  
  
    dict[key] = item + 1
```

Python C-API Demo

```
int
incr_item(PyObject *dict, PyObject *key)
{
    /* Objects all initialized to NULL for Py_XDECREF */
    PyObject *item = NULL,
    *const_one = NULL,
    *incremented_item = NULL;
    /* Return value initialized to -1 (failure) */
    int rv = -1;

    item = PyObject_GetItem(dict, key);
    if (item == NULL) {
        /* Handle KeyError only: */
        if (!PyErr_ExceptionMatches(PyExc_KeyError))
            goto error;
    }
}
```

```

    /* Clear the error and use zero: */
    PyErr_Clear();
    item = PyInt_FromLong(0L);
    if (item == NULL)
        goto error;
}

const_one = PyInt_FromLong(1L);
if (const_one == NULL)
    goto error;

incremented_item = PyNumber_Add(item, const_one);
if (incremented_item == NULL)
    goto error;

if (PyObject_SetItem(dict, key, incremented_item) < 0)
    goto error;

rv = 0; /* Success */
/* Continue with cleanup code */

```

error:

```
/* Cleanup code, shared by success and failure path */  
  
/* Use Py_XDECREF() to ignore NULL references */  
Py_XDECREF(item);  
Py_XDECREF(const_one);  
Py_XDECREF(incremented_item);  
  
return rv; /* -1 for error, 0 for success */  
}
```

Cython Advantages

- 99% Python
- Python 2/3 compatibility
- Classes
- Garbage collection
- String handling
- Automatic reference counting
- Automatic type casting (Python->C, C->Python)
- Portable C code produced
- Stable, mature

Cython Disadvantages

- Needs compilation
- Distutils
- CPython specific

Python demo counter

```
def counter(count):  
    x = 0  
    for i in xrange(count):  # range in Py3  
        x += i
```

Cython demo counter

```
def counter(count):  
    cdef int x = 0 # <- a C style data type  
    for i in xrange(count):  
        x += i
```

Cython cdef-ed demo counter

```
cdef int counter(int count):  
    cdef int x = 0  
    for i in xrange(count):  
        x += i  
    return x
```

Building a Cython module

- **Cython translates from .pyx to C code**

cython inputfile.pyx

- **Or let setup.py handle it**

python setup.py build_ext --inplace

Cython and the GIL



Bypassing the GIL with C

Modules that release the GIL:

- `time.sleep()`
- most of NumPy
- many C extensions

Cython nogil

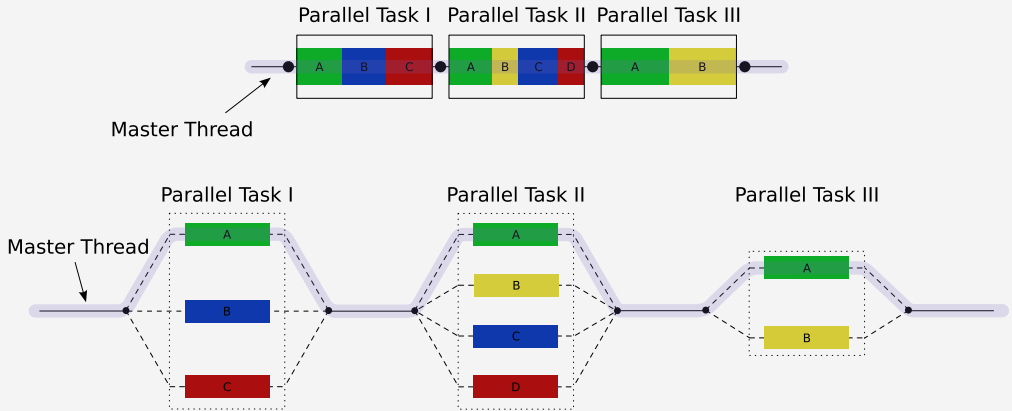
```
def cython_func():  
    with nogil:  
        do_something()  
  
    if something_bad == True:  
        with gil:  
            raise RuntimeError('sorry...')
```


Threading headaches:

- race conditions
- deadlocks
- data corruption
- thread pools
- Yikes! Wait...

Easier multithreading... from C?!

OpenMP: Shared memory multithreading C API/spec



source: <https://en.wikipedia.org/wiki/OpenMP>

Classic Demo Updated

- **2D Laplace Equation benchmark by Prabhu Ramachandran in 2004:**

<http://wiki.scipy.org/PerformancePython>

- **Updated in by Travis Oliphant in 2011:**

<http://technicaldiscovery.blogspot.co.nz/2011/06/speeding-up-python-numpy-cy>

- **Previously compared:**

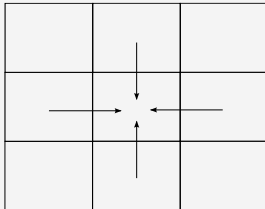
Psyco, NumPy, Blitz, Inline, Python/Fortran, Pyrex, MatLab, Octave, Pure C++

- **We'll discuss:**

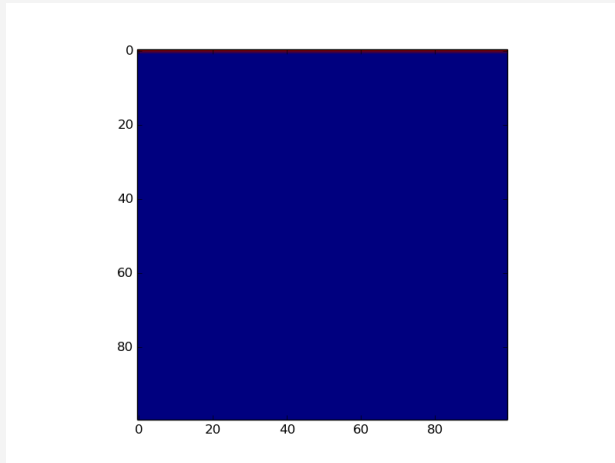
Python, NumPy, Numba, Cython, Cython wrapping C, Cython in parallel

2D Laplace equation

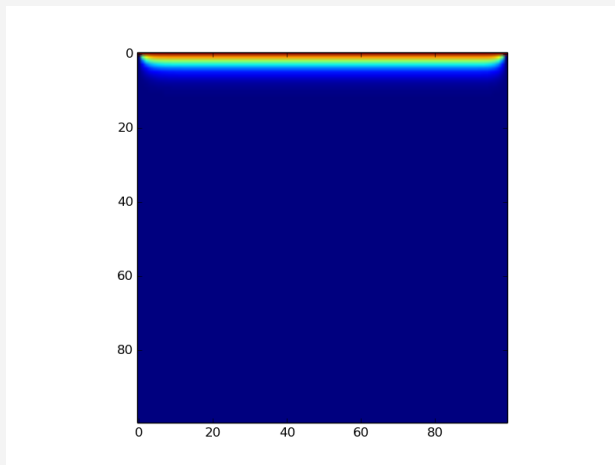
- floating point intensive
- iterative



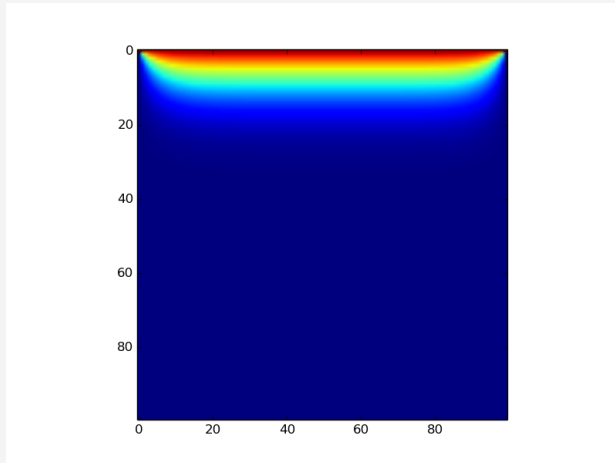
Starting state



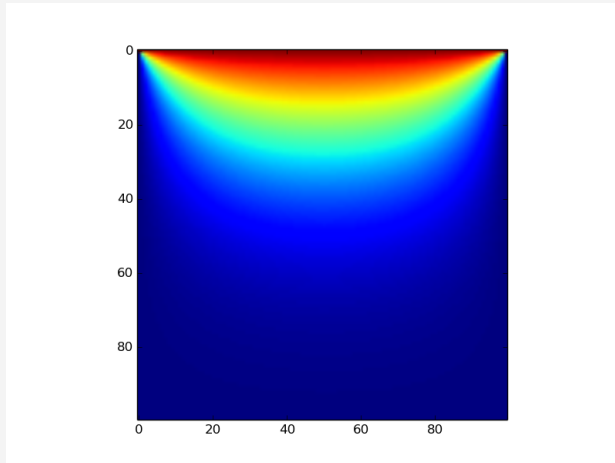
10 iterations




100 iterations



1000 iterations



10000 iterations




Python version

```
def py_update(u, dx2, dy2):
    nx, ny = u.shape
    for i in xrange(1,nx-1):
        for j in xrange(1, ny-1):
            u[i,j] = ((u[i+1, j] + u[i-1, j]) * dy2 +
                      (u[i, j+1] + u[i, j-1]) * dx2) / (2*(dx2+dy2))

work_array = np.zeros([array_shape, array_shape],
                      dtype=np.float64)
work_array[0] = 1.0

for x in range(100):
    py_update(work_array, dx2, dy2)
```

Python benchmark





Numpy version

- Eliminates all loops
- Extensive use of NumPy vectorized operations
- Creates several temporary arrays

```
import numpy as np

def num_update(u, dx2, dy2):
    u[1:-1,1:-1] = ((u[2:,1:-1] + u[:-2,1:-1])*dy2 +
                    (u[1:-1,2:] + u[1:-1,:-2])*dx2)/(2*(dx2+dy2))
```

Numpy Benchmark




Numba version

- Identical to Python version apart from jit decorator

```
from numba import jit

@jit
def numba_update(u, dx2, dy2):
    for i in xrange(1,u.shape[0]-1):
        for j in xrange(1, u.shape[1]-1):
            u[i,j] = ((u[i+1, j] + u[i-1, j]) * dy2 +
                      (u[i, j+1] + u[i, j-1]) * dx2) / (2*(dx2+dy2))
```

Numba benchmark



Cython version

- Similar to the Python and Numba versions

```
import numpy as np
cimport numpy as np
cimport cython

@cython.boundscheck(False)
@cython.wraparound(False)
def cy_update(np.ndarray[double, ndim=2] u,
              double dx2,
              double dy2):
    cdef int i, j
    for i in xrange(1,u.shape[0]-1):
        for j in xrange(1, u.shape[1]-1):
            u[i,j] = ((u[i+1, j] + u[i-1, j]) * dy2 +
                      (u[i, j+1] + u[i, j-1]) * dx2) / (2*(dx2+dy2))
```


Cython version: setup.py

```
from distutils.core import setup
from distutils.extension import Extension
from Cython.Build import cythonize

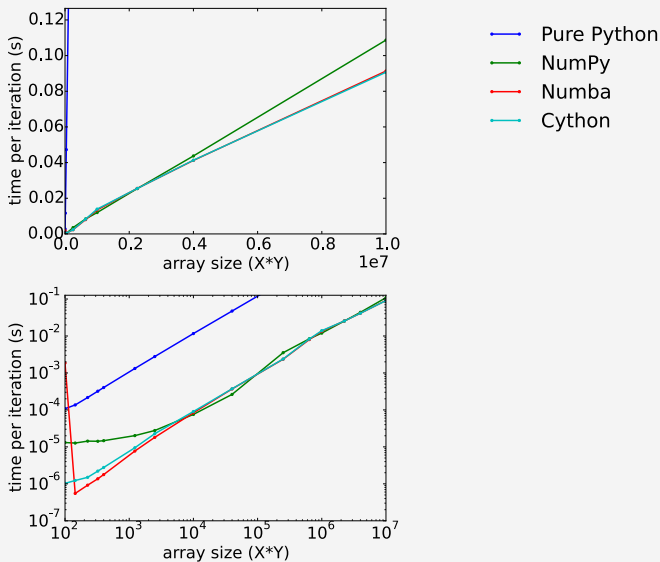
extensions = [Extension('cy_laplace', ['cy_laplace.pyx'])]

setup(name = 'Demos', ext_modules = cythonize(extensions))
```

Build module with a single command:

```
python setup.py build_ext --inplace
```

Cython benchmark



Cython C wrapper

- Calls a C Laplace implementation. Similar to Ctypes or CFFI.

```
import numpy as np
cimport numpy as np

cdef extern from "claplace.h":
    void c_update(double *u, int x_len, int y_len,
                  double dx2, double dy2)

def cy_update_c_wrap(np.ndarray[double, ndim=2] u, dx2, dy2):
    c_update(<double *> &u[0,0], u.shape[0], u.shape[1], dx2, dy2)
```

C implementation

C code in a Python talk?!

```
void c_update(double *u,
              int nx,
              int ny,
              double dx2,
              double dy2) {
    int i, j, idx;
    for (i=1; i<ny-1; i++) {
        for (j=1; j<nx-1; j++) {
            idx = i*nx + j;
            u[idx] = ((u[idx+nx] + u[idx-nx]) * dy2 +
                      (u[idx+1] + u[idx-1]) * dx2) / (2*(dx2+dy2));
        }
    }
}
```

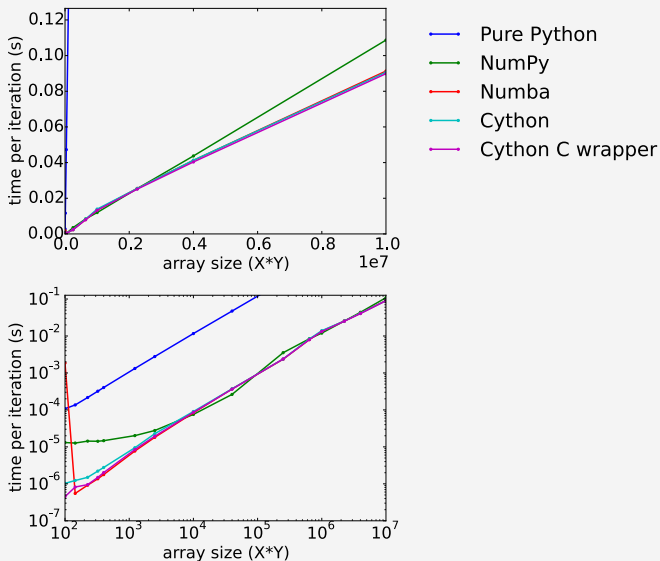
Cython C wrapper: setup.py

```
from distutils.core import setup
from distutils.extension import Extension
from Cython.Build import cythonize

extensions = [Extension('cy_wrap_claplace',
                        ['cy_wrap_claplace.pyx',
                         'claplace.c'],
                        #extra_compile_args=['-fopenmp'],
                        #extra_link_args=['-fopenmp']
                        )
              ]

setup(name = 'Demos', ext_modules = cythonize(extensions))
```

Cython C wrapper benchmark



Cython parallelism

Cython.parallel

- parallel directive: thread-local buffers
- prange: an OpenMP parallel version of xrange/range

openmp module

- low level C API wrappers, eg:


```
num_threads = openmp.omp_get_num_threads()
```

Cython parallel version

```
@cython.boundscheck(False)
@cython.wraparound(False)
def cy_update_parallel(np.ndarray[double, ndim=2] u,
                       double dx2,
                       double dy2):

    cdef int i, j
    for i in prange(1, u.shape[0]-1, nogil=True):
        for j in xrange(1, u.shape[1]-1):
            u[i,j] = ((u[i+1, j] + u[i-1, j]) * dy2 +
                      u[i, j+1] + u[i, j-1]) * dx2) / (2*(dx2+dy2))
```


Cython parallel benchmark



Need more performance?

- Compiler flags (-O3, -ffast-math, ...)
- PyCuda/PyOpenCL
- NumbaPro
- OpenMP 4, OpenACC
- Distributed parallelism: mpi4py, ipython parallel, Spark, etc

Conclusions

- Cython make C extensions easy
- Excellent performance, especially in parallel
- Numba also impressive, but no prange

Arbitrary scores:

	Pure Python	NumPy	Cython	Cython parallel	Numba
Simplicity	***	****	**	**	***
Performance		***	****	*****	****
Distribution	*****	****	**	**	*

Thanks!

- <http://cython.org>
- <http://numba.pydata.org>
- https://github.com/crleblanc/cython_talk_2105

Questions?